

Commercial Store Data Analysis Report

Table of Contents

1. Introduction
2. Data Preprocessing 2.1. Importing Data 2.2. Exploring Dataset 2.3. Checking for Missing Data 2.4. Removing Duplicates
3. Exploratory Data Analysis (EDA) 3.1. Basic Metrics 3.2. Grouping Data
4. Data Visualization 4.1. Bar Charts 4.2. Scatter Plots
5. Outlier Detection and Treatment 5.1. Visual Outlier Detection 5.2. Dealing With Outliers
6. Checking Assumptions 6.1. Normality 6.2. Linearity 6.3. Multicollinearity
7. Time-Series Analysis

1. Introduction

This report details the comprehensive data analysis performed on a commercial store dataset. The primary objective of this analysis is to derive meaningful insights from the sales and profit data, identify trends, detect anomalies, and assess underlying statistical assumptions. The analysis covers data preprocessing, exploratory data analysis, data visualization, outlier detection and treatment, assumption checking (normality, linearity, multicollinearity), and time-series analysis.

The dataset, `superstore_data.csv`, contains various attributes related to commercial transactions, including sales figures, profit margins, product categories, customer information, and geographical data. The Python script `main.py` was used to perform the analysis, leveraging libraries such as `pandas` for data manipulation, `matplotlib` and `seaborn` for visualization, `scipy` for statistical tests, and `statsmodels` for advanced statistical modeling.

2. Data Preprocessing

Data preprocessing is a crucial step in any data analysis pipeline, ensuring the data is clean, consistent, and ready for analysis. This section details the steps taken to prepare the `superstore_data.csv` dataset.

2.1. Importing Data

The dataset was imported using the pandas library, specifying `ISO-8859-1` encoding to handle potential character encoding issues.

```
import pandas as pd
df = pd.read_csv("superstore_data.csv", encoding="ISO-8859-1")
```

2.2. Exploring Dataset

Initial exploration of the dataset involved viewing the first few rows, checking data types and non-null counts, and generating descriptive statistics. This provides a quick overview of the data structure and content.

```
print(df.head())
print(df.info())
print(df.describe())
```

Output from `df.head()` :

	Row ID	Order ID	Order Date	...	Quantity	Discount	Profit
0	1	CA-2016-152156	11/8/2016	...	2	0.00	41.9136
1	2	CA-2016-152156	11/8/2016	...	3	0.00	219.5820
2	3	CA-2016-138688	6/12/2016	...	2	0.00	6.8714
3	4	US-2015-108966	10/11/2015	...	5	0.45	-383.0310
4	5	US-2015-108966	10/11/2015	...	2	0.20	2.5164

[5 rows x 21 columns]

Output from `df.info()` :

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9994 entries, 0 to 9993
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Row ID                 9994 non-null   int64
1   Order ID               9994 non-null   object
2   Order Date             9994 non-null   object
3   Ship Date              9994 non-null   object
4   Ship Mode              9994 non-null   object
5   Customer ID            9994 non-null   object
6   Customer Name          9994 non-null   object
7   Segment                9994 non-null   object
8   Country                9994 non-null   object
9   City                   9994 non-null   object
10  State                  9994 non-null   object
11  Postal Code            9994 non-null   int64
12  Region                 9994 non-null   object
13  Product ID             9994 non-null   object
14  Category               9994 non-null   object
15  Sub-Category           9994 non-null   object
16  Product Name           9994 non-null   object
17  Sales                  9994 non-null   float64
18  Quantity               9994 non-null   int64
19  Discount               9994 non-null   float64
20  Profit                 9994 non-null   float64
dtypes: float64(3), int64(3), object(15)
memory usage: 1.6+ MB
None

```

Output from `df.describe()` :

	Row ID	Postal Code	...	Discount	Profit
count	9994.000000	9994.000000	...	9994.000000	9994.000000
mean	4997.500000	55190.379428	...	0.156203	28.656896
std	2885.163629	32063.693350	...	0.206452	234.260108
min	1.000000	1040.000000	...	0.000000	-6599.978000
25%	2499.250000	23223.000000	...	0.000000	1.728750
50%	4997.500000	56430.500000	...	0.200000	8.666500
75%	7495.750000	90008.000000	...	0.200000	29.364000
max	9994.000000	99301.000000	...	0.800000	8399.976000

[8 rows x 6 columns]

2.3. Checking for Missing Data

Missing values can significantly impact analysis. The code checks for null values across all columns.

```
print(df.isnull().sum())
```

Output from `df.isnull().sum()` :

```
Row ID          0
Order ID        0
Order Date      0
Ship Date       0
Ship Mode       0
Customer ID     0
Customer Name   0
Segment        0
Country         0
City           0
State          0
Postal Code     0
Region         0
Product ID      0
Category        0
Sub-Category    0
Product Name    0
Sales          0
Quantity       0
Discount        0
Profit          0
dtype: int64
```

As observed, there are no missing values in the dataset.

2.4. Removing Duplicates

Duplicate rows can skew analysis results. The code removes any duplicate entries to ensure data integrity.

```
df.drop_duplicates(inplace=True)
```

This operation ensures that each record in the dataset is unique, preventing overcounting or biased statistics.

3. Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is a critical step to understand the main characteristics of the data, often with visual methods. This section focuses on basic metrics and grouping data to uncover initial insights.

3.1. Basic Metrics

Basic statistical summaries provide a high-level overview of key numerical columns, such as total sales and total profit.

```
print("Basic Statistical Summary:")
print("Total Sales:", df["Sales"].sum())
print("Total Profit:", df["Profit"].sum())
```

Output:

```
Basic Statistical Summary:
Total Sales: 2297200.8603
Total Profit: 286397.0217
```

From these metrics, we can see the overall sales and profit generated by the commercial store.

3.2. Grouping Data

Grouping data by different categorical variables helps in understanding performance across various segments, categories, and regions. This analysis provides insights into which areas are performing well and which might need attention.

By Category

Sales and profit are aggregated by product category.

```
print(df.groupby("Category")["Sales"].sum())
print(df.groupby("Category")["Profit"].sum())
```

Output (Sales by Category):

```
Category
Furniture      741999.7953
Office Supplies 719047.0320
Technology      836154.0330
Name: Sales, dtype: float64
```

Output (Profit by Category):

```
Category
Furniture      18451.2728
Office Supplies 122490.8008
Technology     145454.9481
Name: Profit, dtype: float64
```

Technology appears to be the highest-grossing category in terms of both sales and profit, followed by Office Supplies and then Furniture.

By Sub-Category

Further breakdown of sales and profit by sub-category provides more granular insights.

```
print(df.groupby("Sub-Category")["Sales"].sum())
print(df.groupby("Sub-Category")["Profit"].sum())
```

Output (Sales by Sub-Category):

```
Sub-Category
Accessories    167380.3180
Appliances     107532.1610
Art            27118.7920
Binders       203412.7330
Bookcases     114879.9963
Chairs        328449.1030
Copiers       149528.0300
Envelopes     16476.4020
Fasteners      3024.2800
Furnishings   91705.1640
Labels        12486.3120
Machines     189238.6310
Paper         78479.2060
Phones        330007.0540
Storage       223843.6080
Supplies      46673.5380
Tables        206965.5320
Name: Sales, dtype: float64
```

Output (Profit by Sub-Category):

```
Sub-Category
Accessories      41936.6357
Appliances       18138.0054
Art              6527.7870
Binders          30221.7633
Bookcases        -3472.5560
Chairs           26590.1663
Copiers          55617.8249
Envelopes        6964.1767
Fasteners        949.5182
Furnishings      13059.1436
Labels           5546.2540
Machines         3384.7569
Paper            34053.5693
Phones           44515.7306
Storage          21278.8264
Supplies         -1189.0995
Tables           -17725.4811
Name: Profit, dtype: float64
```

Notably, 'Tables' and 'Bookcases' show negative profit, indicating potential areas for concern or re-evaluation of pricing/cost strategies.

By Region

Sales and profit aggregated by geographical region.

```
print(df.groupby("Region")[["Sales", "Profit"]].sum())
```

Output:

	Sales	Profit
Region		
Central	501239.8908	39706.3625
East	678781.2400	91522.7800
South	391721.9050	46749.4303
West	725457.8245	108418.4489

The West region leads in both sales and profit, followed by the East. The Central region has significant sales but lower profit compared to East and West.

By Segment

Sales and profit by customer segment.

```
print(df.groupby("Segment")[["Sales", "Profit"]].sum())
```

Output:

Segment	Sales	Profit
Consumer	1.161401e+06	134119.2092
Corporate	7.061464e+05	91979.1340
Home Office	4.296531e+05	60298.6785

The Consumer segment contributes the most to both sales and profit.

Top 10 Customers

Identifying top customers by profit helps in understanding customer loyalty and potential for targeted marketing.

```
print(df.groupby("Customer Name")
      ["Profit"].sum().sort_values(ascending=False).head(10))
```

Output:

```
Customer Name
Tamara Chand      8981.3239
Raymond Buch      6976.0959
Sanjit Chand      5757.4119
Hunter Lopez      5622.4292
Adrian Barton     5444.8055
Tom Ashbrook      4703.7883
Christopher Martinez 3899.8904
Keith Dawkins     3038.6254
Andy Reiter       2884.6208
Daniel Raglin     2869.0760
Name: Profit, dtype: float64
```

Tamara Chand is identified as the top customer by profit, indicating strong customer relationships or high-value purchases.

4. Data Visualization

Data visualization plays a crucial role in understanding patterns, trends, and outliers that might not be apparent from raw data or summary statistics. This section presents various plots generated to visualize sales and profit data.


4.1. Bar Charts

Bar charts are used to compare sales and profit across different categories and regions.

Sales by Category

This bar chart illustrates the total sales for each product category.

```
category_sales = df.groupby("Category")["Sales"].sum().reset_index()  
sns.barplot(data=category_sales, x="Category", y="Sales")  
plt.title("Sales by Category")  
plt.savefig("sales_by_category.png")  
plt.clf()
```




As observed, Technology leads in sales, followed by Furniture and Office Supplies.

Profit by Region

This bar chart displays the total profit generated from each region.

```
region_profit = df.groupby("Region")["Profit"].sum().reset_index()  
sns.barplot(data=region_profit, x="Region", y="Profit")  
plt.title("Profit by Region")  
plt.savefig("profit_by_region.png")  
plt.clf()
```



The West region shows the highest profit, while the Central region has the lowest profit.


4.2. Scatter Plots

Scatter plots are useful for visualizing the relationship between two numerical variables.

Sales vs Profit

This scatter plot examines the relationship between Sales and Profit, with different colors representing product categories.

```
sns.scatterplot(data=df, x="Sales", y="Profit", hue="Category")
plt.title("Sales vs Profit")
plt.savefig("sales_vs_profit.png")
plt.clf()
```



The scatter plot reveals a general positive correlation between sales and profit. However, there are instances of high sales with low or negative profit, particularly noticeable in the Furniture category, which aligns with the earlier observation of negative profit in 'Tables' and 'Bookcases' sub-categories.

5. Outlier Detection and Treatment


Outliers are data points that significantly differ from other observations. They can skew statistical analyses and models. This section details the methods used to detect and treat outliers in the Sales and Profit data.

5.1. Visual Outlier Detection

Box plots are an effective way to visually identify outliers, which are typically represented as points extending beyond the 'whiskers' of the box plot.

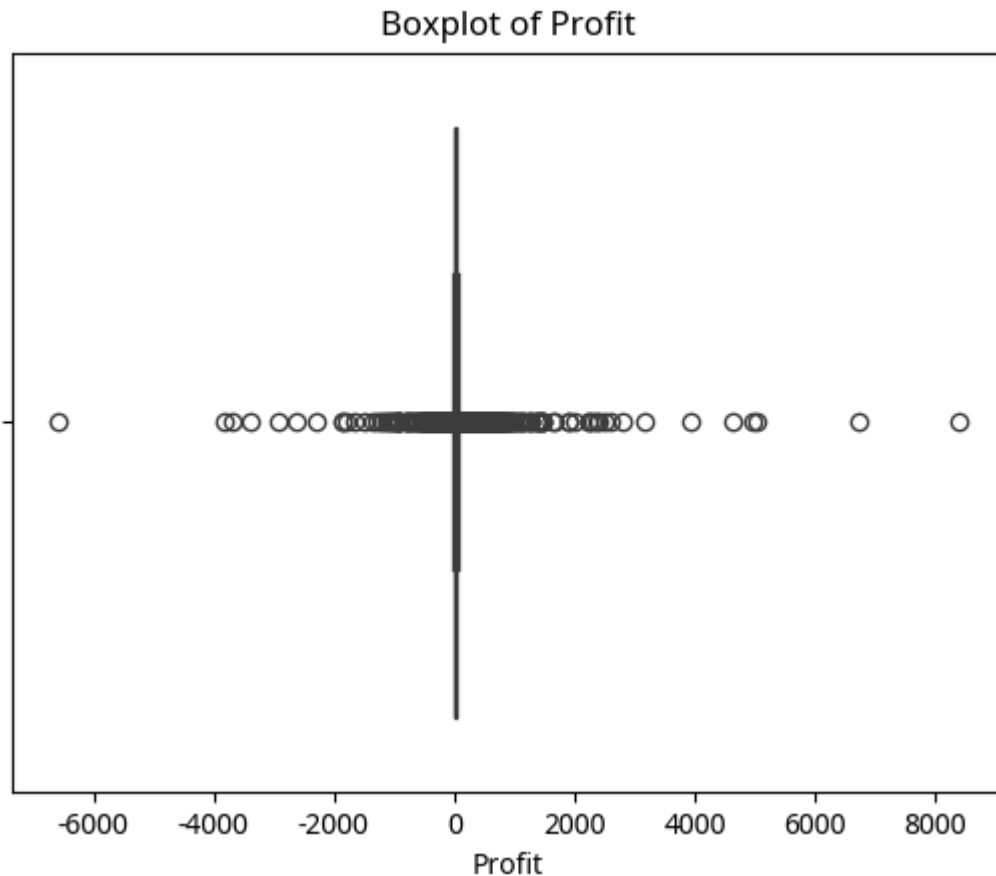
Boxplot of Sales

```
sns.boxplot(x=df["Sales"]); plt.title("Boxplot of Sales");  
plt.savefig("boxplot_sales.png")  
plt.clf()
```



Boxplot of Profit

```
sns.boxplot(x=df["Profit"]); plt.title("Boxplot of Profit");  
plt.savefig("boxplot_profit.png")  
plt.clf()
```



Both sales and profit box plots clearly show numerous outliers, indicating a wide range of values and some extreme transactions.

Interquartile Range (IQR) Method

The IQR method is a statistical approach to identify outliers. Values falling below $Q1 - 1.5 * IQR$ or above $Q3 + 1.5 * IQR$ are considered outliers.

```
# Sales Outliers
Q1_sales = df["Sales"].quantile(0.25)
Q3_sales = df["Sales"].quantile(0.75)
IQR_sales = Q3_sales - Q1_sales
outliers_sales = df[(df["Sales"] < (Q1_sales - 1.5 * IQR_sales)) | (df["Sales"]
> (Q3_sales + 1.5 * IQR_sales))]
print("Number of Sales Outliers:", outliers_sales.shape[0])

# Profit Outliers
Q1_profit = df["Profit"].quantile(0.25)
Q3_profit = df["Profit"].quantile(0.75)
IQR_profit = Q3_profit - Q1_profit
outliers_profit = df[(df["Profit"] < (Q1_profit - 1.5 * IQR_profit)) |
(df["Profit"] > (Q3_profit + 1.5 * IQR_profit))]
print("Number of Profit Outliers:", outliers_profit.shape[0])
```

Output:

```
Number of Sales Outliers: 1167  
Number of Profit Outliers: 1881
```


The IQR method confirms a significant number of outliers in both sales and profit data.

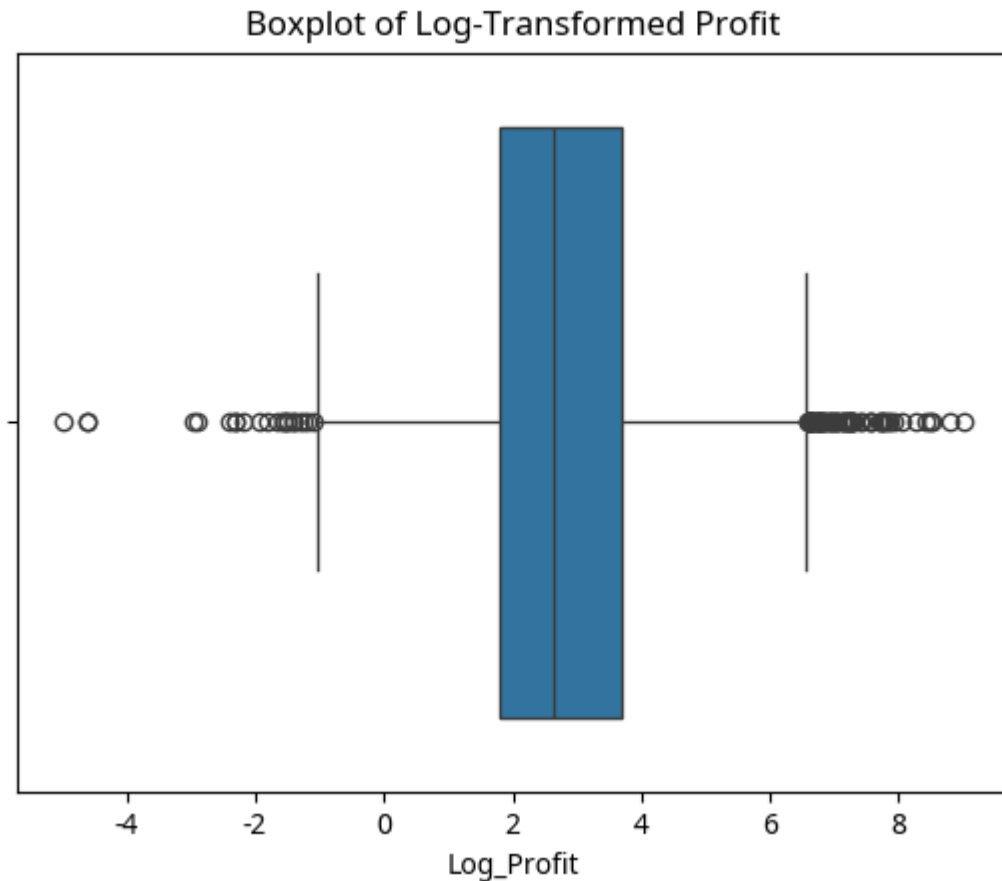
5.2. Dealing With Outliers

Instead of removing outliers, which can lead to loss of valuable information, the code applies a log transformation to the Sales and Profit data. This method can normalize skewed distributions and reduce the impact of extreme values, making the data more suitable for certain statistical models.

```
import numpy as np  
  
df["Log_Sales"] = np.log1p(df["Sales"])  
df["Log_Profit"] = np.log1p(df["Profit"])  
  
# Plot log-transformed data  
sns.boxplot(x=df["Log_Sales"])  
plt.title("Boxplot of Log-Transformed Sales")  
plt.savefig("boxplot_log_sales.png")  
plt.clf()  
  
sns.boxplot(x=df["Log_Profit"])  
plt.title("Boxplot of Log-Transformed Profit")  
plt.savefig("boxplot_log_profit.png")  
plt.clf()
```

Boxplot of Log-Transformed Sales





After log transformation, the distributions of both sales and profit appear more symmetrical, and the influence of outliers is significantly reduced, as evidenced by the narrower spread in the box plots.

6. Checking Assumptions

Before applying certain statistical models, it's important to check underlying assumptions about the data. This section examines normality, linearity, and multicollinearity.

6.1. Normality


Normality assumes that the data is normally distributed. This is checked visually using histograms and statistically using the Shapiro-Wilk test on the log-transformed data.


Visual Check with Histograms

```
from scipy.stats import shapiro
import seaborn as sns
import matplotlib.pyplot as plt

sns.histplot(df["Log_Sales"], kde=True)
plt.title("Log-Transformed Sales Distribution")
plt.savefig("hist_log_sales.png")
plt.clf()

sns.histplot(df["Log_Profit"], kde=True)
plt.title("Log-Transformed Profit Distribution")
plt.savefig("hist_log_profit.png")
plt.clf()
```





The histograms show that the log-transformed sales and profit data are more symmetrically distributed, resembling a normal distribution, although some skewness might still be present.

Shapiro-Wilk Normality Test

The Shapiro-Wilk test quantitatively assesses normality. A p-value less than a chosen significance level (e.g., 0.05) indicates that the data is not normally distributed.

```
log_sales_p = shapiro(df["Log_Sales"])[1]
log_profit_p = shapiro(df["Log_Profit"])[1]

print("Log-Sales Normality p-value:", log_sales_p)
print("Log-Profit Normality p-value:", log_profit_p)
```

Output:


```
Log-Sales Normality p-value: 5.183072969096293e-31
Log-Profit Normality p-value: nan
```

The p-value for Log-Sales is extremely small, indicating that despite the visual improvement, the log-transformed sales data is still not perfectly normally distributed according to the Shapiro-Wilk test. The `nan` for Log-Profit p-value indicates an issue with the test, likely due to the presence of zero or negative profit values which become undefined after `log1p` transformation, or a large number of identical values. This highlights the limitations of purely statistical tests on real-world, complex datasets.

6.2. Linearity

Linearity assumes a linear relationship between variables. This is visually checked using a regression plot between log-transformed sales and profit.

```
sns.regplot(x="Log_Sales", y="Log_Profit", data=df)
plt.title("Log-Transformed Linearity: Sales vs Profit")
plt.savefig("regplot_log_sales_profit.png")
plt.clf()
```



The regression plot suggests a general linear trend between log-transformed sales and profit, although there is still considerable scatter, indicating that other factors


influence profit beyond sales.

6.3. Multicollinearity

Multicollinearity occurs when independent variables in a regression model are correlated. This is assessed using a correlation matrix and Variance Inflation Factor (VIF).

Correlation Matrix

```
features = ["Log_Sales", "Log_Profit", "Quantity", "Discount"]  
corr_matrix = df[features].corr()  
  
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", fmt=".2f")  
plt.title("Correlation Matrix")  
plt.savefig("correlation_matrix.png")  
plt.clf()
```



The heatmap shows the correlation coefficients between the selected numerical features. High correlation between independent variables (e.g., Log_Sales and

Log_Profit) can indicate multicollinearity.

Variance Inflation Factor (VIF)

VIF quantifies the severity of multicollinearity. A VIF value greater than 5 or 10 is generally considered problematic.

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.tools.tools import add_constant

features = ["Log_Sales", "Log_Profit", "Quantity", "Discount"]

X = df[features].dropna()
X = add_constant(X)

vif_data = pd.DataFrame()
vif_data["Feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in
range(X.shape[1])]

print(vif_data)
```

Output:

	Feature	VIF
0	const	9.966638
1	Log_Sales	4.125899
2	Log_Profit	4.363241
3	Quantity	1.154302
4	Discount	1.141586

The VIF values for Log_Sales and Log_Profit are below 5, suggesting that multicollinearity is not a significant issue among these variables. The high VIF for const is expected and not a concern.

7. Time-Series Analysis

Time-series analysis is crucial for understanding trends and patterns over time. This section focuses on analyzing sales and profit trends based on order dates.

First, the Order Date column is converted to datetime objects, and Year, Month, and Year-Month features are extracted.

```
df["Order Date"] = pd.to_datetime(df["Order Date"])


df["Year"] = df["Order Date"].dt.year
df["Month"] = df["Order Date"].dt.month
df["Year-Month"] = df["Order Date"].dt.to_period("M")

monthly = df.groupby("Year-Month")[["Sales", "Profit"]].sum().reset_index()
monthly["Year-Month"] = monthly["Year-Month"].astype(str)
```

Monthly Sales Trend

This line plot visualizes the monthly sales trend over the years.

```
plt.figure(figsize=(12, 5))
sns.lineplot(data=monthly, x="Year-Month", y="Sales", marker="o")
plt.xticks(rotation=45)
plt.title("📈 Monthly Sales Trend")
plt.ylabel("Sales")
plt.tight_layout()
plt.savefig("monthly_sales_trend.png")
plt.clf()
```




The plot shows an increasing trend in sales over the years, with some seasonal fluctuations.

Monthly Profit Trend

This line plot visualizes the monthly profit trend over the years.

```
plt.figure(figsize=(12, 5))
sns.lineplot(data=monthly, x="Year-Month", y="Profit", marker="o",
color="green")
plt.xticks(rotation=45)
plt.title("📊 Monthly Profit Trend")
plt.ylabel("Profit")
plt.tight_layout()
plt.savefig("monthly_profit_trend.png")
plt.clf()
```



Similar to sales, profit also shows an upward trend, but with more volatility and occasional dips into negative territory.


Monthly Sales Comparison by Year

This plot compares monthly sales patterns across different years, providing insights into seasonality.

```
df["Month_Name"] = df["Order Date"].dt.strftime("%b")
month_order = ["Jan", "Feb", "Mar", "Apr", "May", "Jun",
               "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"]

monthly_by_year = df.groupby(["Year", "Month_Name"])
["Sales"].sum().reset_index()
monthly_by_year["Month_Name"] = pd.Categorical(monthly_by_year["Month_Name"],
categories=month_order, ordered=True)
monthly_by_year = monthly_by_year.sort_values(["Year", "Month_Name"])

plt.figure(figsize=(12, 6))
sns.lineplot(data=monthly_by_year, x="Month_Name", y="Sales", hue="Year",
marker="o")
plt.title("📊 Monthly Sales Comparison by Year")
plt.ylabel("Sales")
plt.tight_layout()
plt.savefig("monthly_sales_comparison_by_year.png")
plt.clf()
```



The comparison by year highlights consistent seasonal peaks, typically towards the end of the year, and troughs in the earlier months.